

```
import numpy as np
```

```
from .layers import * from .layer_utils import *
```

```
""" This code was originally written for CS 231n at Stanford University (cs231n.stanford.edu). It has been modified in various areas for use in the ECE 239AS class at UCLA. This includes the descriptions of what code to implement as well as some slight potential changes in variable names to be consistent with class nomenclature. We thank Justin Johnson & Serena Yeung for permission to use this code. To see the original version, please visit cs231n.stanford.edu.
```

```
"""
```

```
class TwoLayerNet(object): """ A two-layer fully-connected neural network with ReLU nonlinearity and softmax loss that uses a modular layer design. We assume an input dimension of D, a hidden dimension of H, and perform classification over C classes.
```

The architecture should be affine - relu - affine - softmax.

Note that this class does not implement gradient descent; instead, it will interact with a separate Solver object that is responsible for running optimization.

The learnable parameters of the model are stored in the dictionary self.params that maps parameter names to numpy arrays. """

```
def init(self, input_dim=33232, hidden_dims=100, num_classes=10,
```

```

        dropout=0, weight_scale=1e-3, reg=0.0):
"""
Initialize a new network.

Inputs:
- input_dim: An integer giving the size of the input
- hidden_dims: An integer giving the size of the hidden layer
- num_classes: An integer giving the number of classes to classify
- dropout: Scalar between 0 and 1 giving dropout strength.
- weight_scale: Scalar giving the standard deviation for random
  initialization of the weights.
- reg: Scalar giving L2 regularization strength.
"""
self.params = {}
self.reg = reg

# ===== #
# YOUR CODE HERE:
# Initialize W1, W2, b1, and b2. Store these as self.params['W1'],
# self.params['W2'], self.params['b1'] and self.params['b2']. The
# biases are initialized to zero and the weights are initialized
# so that each parameter has mean 0 and standard deviation weight_scale.
# The dimensions of W1 should be (input_dim, hidden_dim) and the
# dimensions of W2 should be (hidden_dims, num_classes)
# ===== #

self.params['W1'] = np.random.randn(input_dim, hidden_dims) * weight_scale
self.params['b1'] = np.zeros(hidden_dims)
self.params['W2'] = np.random.randn(hidden_dims, num_classes) * weight_scale
self.params['b2'] = np.zeros(num_classes)

# ===== #
# END YOUR CODE HERE
# ===== #

```

def loss(self, X, y=None):

```

"""
Compute loss and gradient for a minibatch of data.

Inputs:
- X: Array of input data of shape (N, d_1, ..., d_k)
- y: Array of labels, of shape (N,). y[i] gives the label for X[i].

Returns:

```

If y is None, then run a test-time forward pass of the model and return:

- scores: Array of shape (N, C) giving classification scores, where scores[i, c] is the classification score for X[i] and class c.

If y is not None, then run a training-time forward and backward pass and return a tuple of:

- loss: Scalar value giving the loss
- grads: Dictionary with the same keys as self.params, mapping parameter names to gradients of the loss with respect to those parameters.

"""

scores = None

=====

YOUR CODE HERE:

Implement the forward pass of the two-layer neural network. Store
the class scores as the variable 'scores'. Be sure to use the layers
you prior implemented.

=====

out, cache = affine_relu_forward(X, self.params['W1'], self.params['b1'])
scores, cache_2 = affine_forward(out, self.params['W2'], self.params['b2'])

=====

END YOUR CODE HERE

=====

If y is None then we are in test mode so just return scores

if y is None:

return scores

loss, grads = 0, {}

=====

YOUR CODE HERE:

Implement the backward pass of the two-layer neural net. Store
the loss as the variable 'loss' and store the gradients in the
'grads' dictionary. For the grads dictionary, grads['W1'] holds
the gradient for W1, grads['b1'] holds the gradient for b1, etc.
i.e., grads[k] holds the gradient for self.params[k].

#

Add L2 regularization, where there is an added cost $0.5 \cdot \text{self.reg} \cdot W^2$
for each W. Be sure to include the 0.5 multiplying factor to
match our implementation.

#

And be sure to use the layers you prior implemented.

=====

data_loss, data_loss_grad = softmax_loss(scores, y)

```

r_loss = .5 * self.reg * (np.sum(self.params['W1'] **2) +
np.sum(self.params['W2'] ** 2))
loss = data_loss + r_loss
# now backwards through the network
# so affine backwards needs the data_loss_grad and the original input and
weights and bias into the last layer.
# that is cache 2 from above
dx_into_next_layer, dw_2, db_2 = affine_backward(data_loss_grad, cache_2)
grads['W2'] = dw_2 + self.reg * self.params['W2'] # add the regularization
derivative
grads['b2'] = db_2

# now we want to use affine_relu_backwards to do everything for us
# its incoming gradient is the dx_into_next_layer
# we need to assemble a tuple (fc_cache, relu_cache) where fc_cache was the
inputs into this fc layer
# and relu_cache was the inputs into relu
# luckily this is just the cache from affine_relu_forward
dx, dw_1, db_1 = affine_relu_backward(dx_into_next_layer, cache)
grads['W1'] = dw_1 + self.reg * self.params['W1']
grads['b1'] = db_1

# ===== #
# END YOUR CODE HERE
# ===== #

return loss, grads

```

class FullyConnectedNet(object): """ A fully-connected neural network with an arbitrary number of hidden layers, ReLU nonlinearities, and a softmax loss function. This will also implement dropout and batch normalization as options. For a network with L layers, the architecture will be

{affine - [batch norm] - relu - [dropout]} x (L - 1) - affine - softmax

where batch normalization and dropout are optional, and the {...} block is repeated L - 1 times.

Similar to the TwoLayerNet above, learnable parameters are stored in the self.params dictionary and will be learned using the Solver class. """

def **init**(self, hidden_dims, input_dim=33232, num_classes=10,

```

        dropout=0, use_batchnorm=False, reg=0.0,
        weight_scale=1e-2, dtype=np.float32, seed=None):
    """
    Initialize a new FullyConnectedNet.

    Inputs:

```

```

- hidden_dims: A list of integers giving the size of each hidden layer.
- input_dim: An integer giving the size of the input.
- num_classes: An integer giving the number of classes to classify.
- dropout: Scalar between 0 and 1 giving dropout strength. If dropout=0 then
  the network should not use dropout at all.
- use_batchnorm: Whether or not the network should use batch normalization.
- reg: Scalar giving L2 regularization strength.
- weight_scale: Scalar giving the standard deviation for random
  initialization of the weights.
- dtype: A numpy datatype object; all computations will be performed using
  this datatype. float32 is faster but less accurate, so you should use
  float64 for numeric gradient checking.
- seed: If not None, then pass this random seed to the dropout layers. This
  will make the dropout layers deterministic so we can gradient check the
  model.
"""
self.use_batchnorm = use_batchnorm
self.use_dropout = dropout > 0
self.reg = reg
self.num_layers = 1 + len(hidden_dims)
self.dtype = dtype
self.params = {}

# ===== #
# YOUR CODE HERE:
#   Initialize all parameters of the network in the self.params dictionary.
#   The weights and biases of layer 1 are W1 and b1; and in general the
#   weights and biases of layer i are Wi and bi. The
#   biases are initialized to zero and the weights are initialized
#   so that each parameter has mean 0 and standard deviation weight_scale.
# ===== #

for idx, val in enumerate(hidden_dims):
    self.params['W' + str(idx+1)] = np.random.randn(input_dim if idx == 0 else
hidden_dims[idx-1],
    val if idx != len(hidden_dims) - 1 else num_classes) * weight_scale
    self.params['b' + str(idx+1)] = np.zeros(val if idx != len(hidden_dims) - 1
else num_classes)
    if self.use_batchnorm and idx != len(hidden_dims) - 1:
        # then we should add batchnorm params
        self.params['gamma' + str(idx + 1)] = np.ones(hidden_dims[idx])
        self.params['beta' + str(idx + 1)] = np.zeros(hidden_dims[idx])

# ===== #
# END YOUR CODE HERE
# ===== #

```

```

# When using dropout we need to pass a dropout_param dictionary to each
# dropout layer so that the layer knows the dropout probability and the mode
# (train / test). You can pass the same dropout_param to each dropout layer.
self.dropout_param = {}
if self.use_dropout:
    self.dropout_param = {'mode': 'train', 'p': dropout}
    if seed is not None:
        self.dropout_param['seed'] = seed

# if use_batchnorm:
#     dims = [input_dim] + hidden_dims + [num_classes]
#     for idx in range(self.num_layers-1):
#         self.params['gamma' + str(idx+1)] = np.ones(dims[idx+1])
#         self.params['beta' + str(idx+1)] = np.zeros(dims[idx+1])
# With batch normalization we need to keep track of running means and
# variances, so we need to pass a special bn_param object to each batch
# normalization layer. You should pass self.bn_params[0] to the forward pass
# of the first batch normalization layer, self.bn_params[1] to the forward
# pass of the second batch normalization layer, etc.
self.bn_params = []
if self.use_batchnorm:
    self.bn_params = [{'mode': 'train'} for i in np.arange(self.num_layers - 1)]

# Cast all parameters to the correct datatype
for k, v in self.params.items():
    self.params[k] = v.astype(dtype)

```

```
def loss(self, X, y=None):
```

```

"""
Compute loss and gradient for the fully-connected net.

Input / output: Same as TwoLayerNet above.
"""
X = X.astype(self.dtype)
mode = 'test' if y is None else 'train'

# Set train/test mode for batchnorm params and dropout param since they
# behave differently during training and testing.
if self.dropout_param is not None:
    self.dropout_param['mode'] = mode
if self.use_batchnorm:
    for bn_param in self.bn_params:
        bn_param[mode] = mode

```

```

scores = None

# ===== #
# YOUR CODE HERE:
# Implement the forward pass of the FC net and store the output
# scores as the variable "scores".
# ===== #
layer_caches = []
input = X
for i in range(1, self.num_layers):
    w, b = self.params['W' + str(i)], self.params['b' + str(i)]
    try:
        gamma, beta = self.params['gamma' + str(i)], self.params['beta' + str(i)]
    except KeyError:
        assert i == self.num_layers - 1 or not self.use_batchnorm, "houston we
have a problem"
        gamma, beta = None, None
    if i == self.num_layers - 1:
        out, cache = affine_forward(out, w, b)

    else:
        if self.use_batchnorm and self.use_dropout:
            # so first, we affine
            out, affine_cache = affine_forward(input, w, b)
            # next, batchnorm
            out, bn_cache = batchnorm_forward(out, gamma, beta, bn_param)
            # next, relu
            out, relu_cache = relu_forward(out)
            # next, dropout
            out, dropout_cache = dropout_forward(out, self.dropout_param)
            cache = (affine_cache, bn_cache, relu_cache, dropout_cache)
        elif self.use_batchnorm and not self.use_dropout:
            # just do those 3 steps above
            # so first, we affine
            out, affine_cache = affine_forward(input, w, b)
            # next, batchnorm
            out, bn_cache = batchnorm_forward(out, gamma, beta, bn_param)
            # next, relu
            out, relu_cache = relu_forward(out)
            cache = (affine_cache, bn_cache, relu_cache)
        elif self.use_dropout and not self.use_batchnorm:
            # just dropout, no batchnorm
            # okay, we're doing dropout. so first do the affine relu forward:
            out, relu_cache = affine_relu_forward(input, w, b)
            # and now do the dropout

```

```

        out, dropout_cache = dropout_forward(out, self.dropout_param)
        cache = (relu_cache, dropout_cache)
    else:
        # no dropout or batchnorm
        out, cache = affine_relu_forward(input, w, b)
    input = out # input into the next layer is the out of this layer
    layer_caches.append(cache) # store each layer inputs basically
scores = out

```

```

# ===== #
# END YOUR CODE HERE
# ===== #

# If test mode return early
if mode == 'test':
    return scores

loss, grads = 0.0, {}
# ===== #
# YOUR CODE HERE:
# Implement the backwards pass of the FC net and store the gradients
# in the grads dict, so that grads[k] is the gradient of self.params[k]
# Be sure your L2 regularization includes a 0.5 factor.
# ===== #

data_loss, data_loss_grad = softmax_loss(scores, y)
r_loss = .5 * self.reg * np.sum([np.sum(self.params['W' + str(i)] **2) for i
in range(1, self.num_layers)])
loss = data_loss + r_loss
# now backprop
incoming_grad = data_loss_grad # the incoming grad into the last layer is the
loss gradient; this will be updated everytime with the newly computed gradient
for i in range(self.num_layers-1, 0, -1):
    w, b = self.params['W' + str(i)], self.params['b' + str(i)]
    # affine_relu_backwards on everything besides the very last layer.
    if i == self.num_layers - 1:
        dx, dw, db = affine_backward(incoming_grad, layer_caches[i-1])
    else:
        if self.use_batchnorm and self.use_dropout:
            # so first we get our caches
            affine_cache, bn_cache, relu_cache, dropout_cache = layer_caches[i-1]
            # now backward
            dx = dropout_backward(incoming_grad, dropout_cache)
            dx = relu_backward(dx, relu_cache)
            dx,dgamma,dbeta = batchnorm_backward(dx, bn_cache)

```



```

        dx, dw, db = affine_backward(dx, affine_cache)
        grads['gamma' + str(i)], grads['beta' + str(i)] = dgamma, dbeta
    elif self.use_batchnorm and not self.use_dropout:
        # get our caches
        affine_cache, bn_cache, relu_cache = layer_caches[i-1]
        dx = relu_backward(incoming_grad, relu_cache)
        dx, dgamma, dbeta = batchnorm_backward(dx, bn_cache)
        dx, dw, db = affine_backward(dx, affine_cache)
        grads['gamma' + str(i)], grads['beta' + str(i)] = dgamma, dbeta
    elif self.use_dropout and not self.use_batchnorm:
        # ok, we were using dropout, so we should break apart our cache.
        relu_backward_cache, dropout_cache = layer_caches[i-1]
        dx = dropout_backward(incoming_grad, dropout_cache)
        dx, dw, db = affine_relu_backward(dx, relu_backward_cache)
    else:
        dx, dw, db = affine_relu_backward(incoming_grad, layer_caches[i-1])
    grads['W' + str(i)], grads['b' + str(i)] = dw + self.reg * w, db
    incoming_grad = dx # assign the gradient signal incoming into the next layer

# ===== #
# END YOUR CODE HERE
# ===== #
return loss, grads

```